

Master 1 Informatique and Master 1 MOSIG

Parallel Algorithms and Programming

Lab #5

April-2019

**Done by
Salman Farhat
Nesrin Mansouri**

How to Run

To run our code you should use:

1. `mpicc -o try PROGX.c -lm`
2. `mpirun -n N try`

N should have an integer square root (9 --> 3, 4-->2)

After you need to enter the matrix Dimension it will be NxN where N should be divisible by the dimension of the process grid n (n X n).

Fox Algorithm

For Part 1 in the project

To write the fox algorithm using MPI we did:

- Create a grid of processes (dim X dim)
- Create communicator on columns and rows
- Create a new datatype MatrixBlockType of (blockSize * blockSize) where blockSize is the dimension of the submatrix of the Matrix
- In each process, we initialize the submatrices
- Then we apply these steps:

```
MPI_Cart_shift(commcol, 0, 1, &dest, &source);
```

Here in each process, we calculate from where to receive and where to send, receive from below send up.

```
for(int stage=0 ; stage < dim ; stage++){
```

Loop on number of columns in process grid which is dim (columns = rows dim*dim)

```
root=(coord_2d[0] + stage ) % dim;
```

We calculate the root for process 0 it will be first 0 then 1 etc ...

```
if(root==coord_2d[1]) {seconds
MPI_Bcast(A, 1, MatrixBlockType, root, commrow);
multiply(A,B,C,blockSize);
}else {
MPI_Bcast(tempa, 1, MatrixBlockType, root, commrow);
multiply(tempa,B,C, blockSize);
```

```
}
```

Here if the root is equal to the the column number it will broadcast it's block A and all others that are on the same row they will receive the A in TempA and proceed multiplication.

```
for(int l=0 ; l < blockSize ; l++){  
    for(int o = 0 ; o <  blockSize ; o++){  
        buff[l*blockSize + o]= B[l*blockSize + o] ;  
    }  
}
```

Here we are coping block B to a buffer to send using Isend

```
MPI_Isend(buff, 1, MatrixBlockType, dest, 1, commcol, &request2);  
MPI_Irecv(buff2, 1, MatrixBlockType, source, 1, commcol, &request);  
MPI_Wait(&request, &status);  
MPI_Wait(&request2, &status);
```

After the process send it's block B to destination we need to receive new block from the source process which is right below and then we wait to have full completion for the send and receive operations.

```
for(int l=0 ; l < blockSize ; l++){  
    for(int o = 0 ; o <  blockSize ; o++){  
        B[l*blockSize + o] = buff2[l*blockSize + o];  
    }  
}
```

After receiving we are putting the received one to B

```
}
```

After this stage, all process will have the correct values of the C Matrix partition

1. Describe the optimizations that you included in your code if any.

We used non-blocking send and receive (MPI_Isend, MPI_Irecv).

We declared our data types

Cannon Algorithm

For Part 2 in the project

To write the fox algorithm using MPI we did:

- Create a grid of processes (dim X dim)
- Create a new datatype MatrixBlockType of (blockSize * blockSize) where blockSize is the dimension of the submatrix of the Matrix
- In each process, we initialize the submatrices
- Then we apply these steps:

```
int LS , RS , US , DS;
```

```
MPI_Cart_shift(cartcomm, 1, coord_2d[0], &LS, &RS);
```

```
MPI_Cart_shift(cartcomm, 0, coord_2d[1], &US, &DS);
```

process at row 0 will not shift left-right

Process on the row 1 will shift by 1

Process on the row 2 will shift by 2

So we are using the row number `coord_2d[0]` to decide how much we need to shift in each process

Same for Up and down w.r.t columns

We copy A to buffA0 and B to BuffB0

After we send and receive in BuffB0 and BuffA0

```
MPI_Isend(buffA0, 1, MatrixBlockType, LS, 1, cartcomm, & requestA0);
```

```
MPI_Irecv(buffA1, 1, MatrixBlockType, RS, 1, cartcomm, & requestA1);
```

```
MPI_Isend(buffB0, 1, MatrixBlockType, US, 1, cartcomm, & requestB0);
```

```
MPI_Irecv(buffB1, 1, MatrixBlockType, DS, 1, cartcomm, & requestB1);
```

```
MPI_Wait( & requestA0, & status);
```

```
MPI_Wait( & requestA1, & status);
```

```
MPI_Wait( & requestB0, & status);
```

```
MPI_Wait( & requestB1, & status);
```

Then we put the received block in BuffB1 in B and BuffA1 in A.

Then process transfer blocks A on left right and B on up and down

Then each process proceed multiplication

```
multiply(A , B , C , blockSize);
```

Now we need to loop from 1 to dim which is (dim X dim) dimension for the process matrix, and do shift by 1 and transfer blocks Between process and proceed multiplication in each process.

```
for (i = 1; i < dim; i++) {
```

```
    MPI_Cart_shift(cartcomm, 1, 1, &LS, &RS);
```

```
    MPI_Cart_shift(cartcomm, 0, 1, &US, &DS);
```

```
    for (int l = 0; l < blockSize; l++) {
```

```
        for (int o = 0; o < blockSize; o++) {
```

```
            buffA0[l * blockSize + o] = A[l * blockSize + o];
```

```
        }
```

```
    }
```

```
    for (int l = 0; l < blockSize; l++) {
```

```
        for (int o = 0; o < blockSize; o++) {
```

```
            buffB0[l * blockSize + o] = B[l * blockSize + o];
```

```
        }
```

```
    }
```

```
    MPI_Isend(buffA0, 1, MatrixBlockType, LS, 1, cartcomm, & requestA0);
```

```
    MPI_Irecv(buffA1, 1, MatrixBlockType, RS, 1, cartcomm, & requestA1);
```

```
    MPI_Isend(buffB0, 1, MatrixBlockType, US, 1, cartcomm, & requestB0);
```

```
    MPI_Irecv(buffB1, 1, MatrixBlockType, DS, 1, cartcomm, & requestB1);
```

```

MPI_Wait( & requestA0, & status);
MPI_Wait( & requestA1, & status);
MPI_Wait( & requestB0, & status);
MPI_Wait( & requestB1, & status);

for (int l = 0; l < blockSize; l++) {
    for (int o = 0; o < blockSize; o++) {
        A[l * blockSize + o] = buffA1[l * blockSize + o];
    }
}
for (int l = 0; l < blockSize; l++) {
    for (int o = 0; o < blockSize; o++) {
        B[l * blockSize + o] = buffB1[l * blockSize + o];
    }
}

multiply(A, B, C, blockSize);
}

```

After this stage, all process will have the correct values of the C Matrix partition.

1. Describe the optimizations that you included in your code if any.

We used non-blocking send and receive (MPI_Isend, MPI_Irecv)

We declared our data types

Data management

For the data management part, we create big matrices in all process and initialize them in process rank 0

```

int *Main_A = NULL,*Main_B = NULL,*Main_C = NULL;
if(rank == 0){
    Main_A = (int *) malloc(matrixSize * matrixSize * sizeof(int));
    Main_B = (int *) malloc(matrixSize * matrixSize * sizeof(int));
    Main_C = (int *) malloc(matrixSize * matrixSize * sizeof(int));
    fillData(Main_A, Main_B, Main_C, matrixSize);
}

```

After we initialize sendcount array that indicates how much elements from matrix each process should take, and displs indicates from where each process should starts.

```

int c = 0;
for(int i = 0 ; i < dim; i++){
    for(int j = 0 ; j < dim; j++){

```

```

    sendcounts[c] = (matrixSize*matrixSize) / numtasks;
    displs[c] = ((i * blockSize)*matrixSize) + j*blockSize;
    c++;
}
}

```

After we need to rearrange the matrix to make the scatter, so this function arranges the array in a way each process take the correct chunk of the matrix.

```

void PrepareForScater(int * Main_A , int *arrangedBuf ,int startFrom ,int strid, int blockSize , int
matrixSize , int numtasks , int *displs){
    int counter = 0;
    for(int i =0 ; i < numtasks ; i++ ){
        startFrom = displs[i];
        //printf("start from is %d " , startFrom);
        for(int j =0 ; j < blockSize ; j++ ){
            for(int k = startFrom ; k < startFrom + blockSize ; k++){
                arrangedBuf[counter] = Main_A[k];
                counter++;
            }
            startFrom = startFrom + strid;
        }
    }
}

```

```

PrepareForScater(Main_A ,arrangedBufA , startFrom , strid , blockSize , matrixSize , numtasks ,
displs);
PrepareForScater(Main_B ,arrangedBufB , startFrom , strid , blockSize , matrixSize , numtasks ,
displs);

```

arrangedBufA is the matrix that is arranged in a right way

Then we call the scatter

```

MPI_Scatter(arrangedBufA ,blockSize * blockSize, MPI_INT , A , 1 , MatrixBlockType , 0,
cartcomm);
MPI_Scatter(arrangedBufB , blockSize * blockSize, MPI_INT , B , 1 , MatrixBlockType , 0,
cartcomm);

```

After we need to gather the C chunks from all process to the main process (rank 0) we do Gather

```

MPI_Gather(C , 1 , MatrixBlockType , Main_C , 1 , MatrixBlockType , 0 , cartcomm);

```

Exercise 4: Performance evaluation

Fox :

Experiment 1:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores.

Number of processes = 4

Matrix size = 4

Result:

Total time it took to do the multiplication is :0.001016

Number of cycles: 2017206 cycles

Experiment 2:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores.

Number of processes = 4

Matrix size = 1000

Result:

Total time it took to do the multiplication is :5.022948

Number of cycles: 3938555105 cycles

Experiment 3:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total we have 64 logical cores

Number of processes = 16

Matrix size = 5000

```

akhorguani@dahu-5:~/parallel_mpi/nisreen$ perf stat -d mpirun -n 16 ./fox
number of tasks is : 16 , dim 4 4
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 4

Fox algo                                3653207778 cycles

Totall time it took to do the multiplication is :77.594628

Performance counter stats for 'mpirun -n 16 ./fox':

    1293893.238896      task-clock (msec)    #    15.654 CPUs utilized
           85,783      context-switches         #    0.066 K/sec
           293         cpu-migrations          #    0.000 K/sec
          284,988      page-faults             #    0.220 K/sec
    3,618,176,093,330   cycles                    #    2.796 GHz              (50.01%)
    5,834,022,225,383   instructions           #    1.61   insn per cycle   (62.51%)
    142,411,469,347    branches                #   110.064 M/sec           (62.51%)
           120,448,197  branch-misses            #    0.08% of all branches   (62.50%)
    2,781,408,931,276   L1-dcache-loads       #  2149.643 M/sec           (62.50%)
    133,847,613,209    L1-dcache-load-misses #    4.81% of all L1-dcache hits (62.51%)
           7,913,029,600 LLC-loads                #    6.116 M/sec           (50.01%)
           4,215,503,958 LLC-load-misses           #   53.27% of all LL-cache hits (50.01%)

    82.653197252 seconds time elapsed

```

Experiment 4:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total we have 64 logical cores.

Number of processes = 25

Matrix size = 5000

```

number of tasks is : 25 , dim 5 5
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 5

Fox algo                                3304559258 cycles

Totall time it took to do the multiplication is :52.827924

Performance counter stats for 'mpirun -n 25 ./fox':

    1506839.895981      task-clock (msec)    #    23.930 CPUs utilized
           310,733      context-switches         #    0.206 K/sec
           2,967         cpu-migrations          #    0.002 K/sec
          321,054      page-faults             #    0.213 K/sec
    4,200,176,213,617   cycles                    #    2.787 GHz              (50.03%)
    6,001,661,329,988   instructions           #    1.43   insn per cycle   (62.52%)
    177,823,221,413    branches                #   118.011 M/sec           (62.50%)
           164,480,999  branch-misses            #    0.09% of all branches   (62.50%)
    2,839,830,165,420   L1-dcache-loads       #  1884.626 M/sec           (62.51%)
    133,499,542,668    L1-dcache-load-misses #    4.70% of all L1-dcache hits (62.50%)
           7,876,326,830 LLC-loads                #    5.227 M/sec           (50.00%)
           2,868,954,430 LLC-load-misses           #   36.43% of all LL-cache hits (50.02%)

    62.969235645 seconds time elapsed

akhorguani@dahu-5:~/parallel_mpi/nisreen$

```


Full Fox :

Experiment 1:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores

Number of processes = 4

Matrix size = 4

Result:

Number of cycles: 9386317 cycles

Total time it took to do the multiplication is: 0.003781

Experiment 2:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores

Number of processes = 4

Matrix size = 1000

Result:

Number of cycles: 487236063 cycles

Total time it took to do the multiplication is: 7.083110

Experiment 3:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total we have 64 logical cores

Number of processes = 16

Matrix size = 5000

```

akhorguani@dahu-5:~/parallel_mpi/nisreen$ perf stat -d mpirun -n 16 ./full_fox

number of tasks is : 16 , dim 4 4
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 4

Fox algo                                2532994040 cycles

Totall time it took to do the multiplication is :81.159986

Performance counter stats for 'mpirun -n 16 ./full_fox':

      1327089.895331      task-clock (msec)    #    15.751 CPUs utilized
           40,522        context-switches               #     0.031 K/sec
            1,125        cpu-migrations                 #     0.001 K/sec
          358,247        page-faults                    #     0.270 K/sec
    3,709,375,276,520      cycles                        #     2.795 GHz              (50.00%)
    5,849,484,042,055      instructions                 #     1.58   insn per cycle   (62.50%)
    145,245,580,522        branches                     #    109.447 M/sec            (62.50%)
        120,251,759        branch-misses                #     0.08% of all branches   (62.50%)
    2,786,531,717,403      L1-dcache-loads              #   2099.731 M/sec            (62.50%)
    134,720,958,085        L1-dcache-load-misses         #     4.83% of all L1-dcache hits (62.50%)
        7,900,834,970      LLC-loads                    #     5.954 M/sec            (50.00%)
        4,093,366,543      LLC-load-misses              #    51.81% of all LL-cache hits (50.00%)

      84.253385774 seconds time elapsed

```

Experiment 4:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total, we have 64 logical cores

Number of processes = 25

Matrix size = 5000

```

akhorguani@dahu-5:~/parallel_mpi/nisreen$ perf stat -d mpirun -n 25 ./full_fox

number of tasks is : 25 , dim 5 5
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 5

Fox algo                                695244998 cycles

Totall time it took to do the multiplication is :63.882631

Performance counter stats for 'mpirun -n 25 ./full_fox':

    1640441.540649      task-clock (msec)    #    24.512 CPUs utilized
           68,787        context-switches               #     0.042 K/sec
            5,000        cpu-migrations                 #     0.003 K/sec
          393,988        page-faults                    #     0.240 K/sec
    4,575,747,210,917      cycles                        #     2.789 GHz              (50.00%)
    6,038,709,384,106      instructions                 #     1.32   insn per cycle   (62.51%)
    185,395,120,102        branches                     #    113.015 M/sec            (62.50%)
        161,173,010        branch-misses                #     0.09% of all branches   (62.50%)
    2,853,256,505,421      L1-dcache-loads              #   1739.322 M/sec            (62.50%)
    133,627,639,621        L1-dcache-load-misses         #     4.68% of all L1-dcache hits (62.50%)
        7,886,059,571      LLC-loads                    #     4.807 M/sec            (50.00%)
        2,865,854,006      LLC-load-misses              #    36.34% of all LL-cache hits (50.01%)

    66.925155705 seconds time elapsed

```

Canon :

Experiment 1:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores.

Number of processes = 4

Matrix size = 4

Number of cycles: 1168420 cycles

Total time it took to do the multiplication is :0.000697

Experiment 2:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores

Number of processes = 4

Matrix size = 1000

Number of cycles: 33340488 cycles

Total time it took to do the multiplication is :3.457254

Experiment 3:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total we have 64 logical cores

Number of processes = 16

Matrix size = 5000

```

akhorguani@dahu-9:~/parallel_mpi/nisreen$ perf stat -d mpirun -n 16 ./cannon
number of tasks is : 16 , dim 4 4
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 4

Fox algo                                2104725784 cycles

Total time it took to do the multiplication is :87.105842

Performance counter stats for 'mpirun -n 16 ./cannon':

   1390011.850311      task-clock (msec)    #    15.274 CPUs utilized
         320,244      context-switches        #    0.230 K/sec
         51,786      cpu-migrations          #    0.037 K/sec
        236,147      page-faults             #    0.170 K/sec
  3,883,897,378,215    cycles                  #    2.794 GHz              (50.01%)
  5,823,717,307,759    instructions            #    1.50   insn per cycle   (62.51%)
  139,624,498,052     branches                 #   100.448 M/sec           (62.50%)
        133,388,661    branch-misses           #    0.10% of all branches   (62.50%)
  2,778,760,347,342    L1-dcache-loads         # 1999.091 M/sec            (62.49%)
  137,019,291,436     L1-dcache-load-misses      #    4.93% of all L1-dcache hits (62.50%)
       7,906,514,956    LLC-loads                #    5.688 M/sec            (50.00%)
       4,760,496,780    LLC-load-misses          #   60.21% of all LL-cache hits (50.01%)

   91.006071295 seconds time elapsed

```

Experiment 4:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total, we have 64 logical cores.

Number of processes = 25

Matrix size = 5000

```

akhorguani@dahu-9:~/parallel_mpi/nisreen$ perf stat -d mpirun -n 25 ./cannon
number of tasks is : 25 , dim 5 5
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 5

Fox algo                                589297308 cycles

Total time it took to do the multiplication is :74.082357

Performance counter stats for 'mpirun -n 25 ./cannon':

   1779244.205761      task-clock (msec)    #    22.897 CPUs utilized
        1,462,392      context-switches        #    0.822 K/sec
         41,394      cpu-migrations          #    0.023 K/sec
        271,874      page-faults             #    0.153 K/sec
  4,964,270,087,826    cycles                  #    2.790 GHz              (49.98%)
  5,902,665,109,669    instructions            #    1.19   insn per cycle   (62.46%)
  155,434,977,098     branches                 #    87.360 M/sec           (62.48%)
        237,096,629    branch-misses           #    0.15% of all branches   (62.50%)
  2,803,577,917,407    L1-dcache-loads         # 1575.713 M/sec            (62.52%)
  134,640,407,486     L1-dcache-load-misses      #    4.80% of all L1-dcache hits (62.54%)
       7,977,253,617    LLC-loads                #    4.484 M/sec            (50.02%)
       4,889,006,402    LLC-load-misses          #   61.29% of all LL-cache hits (49.99%)

   77.707492466 seconds time elapsed

```

Full Canon :

Experiment 1:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores.

Number of processes = 4

Matrix size = 4

Number of cycles: 1227127 cycles

Total time it took to do the multiplication is :0.000626

Experiment 2:

Hardware description: on 4 cores

2 physical cores each having hyper thread (2 threads) so in total we have 4 logical cores.

Number of processes = 4

Matrix size = 1000

Number of cycles: 401592955 cycles

Total time it took to do the multiplication is :3.609431

Experiment 3:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total we have 64 logical cores

Number of processes = 16

Matrix size = 5000

```

akhorguani@dahu-9:~/parallel_mpi/nisreen$ perf stat -d mpirun -n 16 ./full_cannon

number of tasks is : 16 , dim 4 4
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 4

Full Canon algo                      3480526122 cycles

Total time it took to do the multiplication is :81.718370

Performance counter stats for 'mpirun -n 16 ./full_cannon':

    1339869.855618      task-clock (msec)    #    15.728 CPUs utilized
           49,314      context-switches        #     0.037 K/sec
           634         cpu-migrations          #     0.000 K/sec
          382,445       page-faults            #     0.285 K/sec
 3,745,235,587,930      cycles                  #    2.795 GHz              (50.01%)
 5,861,465,198,671      instructions            #    1.57   insn per cycle   (62.51%)
 147,010,607,738        branches                #   109.720 M/sec            (62.50%)
 120,188,855            branch-misses           #    0.08% of all branches   (62.50%)
 2,791,530,362,231      L1-dcache-loads        # 2083.434 M/sec            (62.50%)
 134,926,125,030        L1-dcache-load-misses     #    4.83% of all L1-dcache hits (62.50%)
 7,900,662,696          LLC-loads                #    5.897 M/sec            (50.00%)
 4,009,410,468          LLC-load-misses           #   50.75% of all LL-cache hits (50.01%)

    85.189507000 seconds time elapsed

```

Experiment 4:

Hardware description: on 64 cores

32 physical cores each having hyper thread (2 threads) so in total, we have 64 logical cores

Number of processes = 25

Matrix size = 5000

```

akhorguani@dahu-9:~/parallel_mpi/nisreen$ perf stat -d mpirun -n 25 ./full_cannon

number of tasks is : 25 , dim 5 5
Matrix is NxN it should be N mod dim = 0   enter N:
5000

matrix columns size is : 5000 , process columns size is 5

Full Canon algo                      3509716226 cycles

Total time it took to do the multiplication is :69.433487

Performance counter stats for 'mpirun -n 25 ./full_cannon':

    1755242.768422      task-clock (msec)    #    24.133 CPUs utilized
       249,803          context-switches        #     0.142 K/sec
       12,141           cpu-migrations          #     0.007 K/sec
       418,583          page-faults            #     0.238 K/sec
 4,905,958,779,867      cycles                  #    2.795 GHz              (49.99%)
 6,075,268,625,701      instructions            #    1.24   insn per cycle   (62.49%)
 192,214,327,203        branches                #   109.509 M/sec            (62.50%)
 171,560,790            branch-misses           #    0.09% of all branches   (62.51%)
 2,867,238,403,128      L1-dcache-loads        # 1633.528 M/sec            (62.52%)
 133,958,905,824        L1-dcache-load-misses     #    4.67% of all L1-dcache hits (62.51%)
 7,886,264,776          LLC-loads                #    4.493 M/sec            (50.00%)
 2,524,868,009          LLC-load-misses           #   32.02% of all LL-cache hits (50.00%)

    72.733331309 seconds time elapsed

```

References

“Efficient Parallel Implementation of the Fox Algorithm.” *CiteSeerX*,
citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.84.7509.

An Introduction to Parallel Computing, parallelcomp.uw.hu/ch06lev1sec7.html.

Special Thanks to Ana Khorguani. She helped me in testing phase where she run my algorithm on the 64 cores she has access to.